

Part I: Naive Base Classifier

| Project | F1 |
|------------|---------------------|
| jackrabbit | 0.2686414708886619 |
| xorg | 0.3061667183142237 |
| jdt | 0.2731565564838358 |
| lucene | 0.35405405405405405 |

Part II

| Project | F1/Logistic Regression | F1/Decission Tree | F1/SVM | F1/Random Forest |
|------------|------------------------|---------------------|----------------------|---------------------|
| jackrabbit | 0.38014311270125 | 0.3705035971223 | 0.3539094650205762 | 0.2759336099585062 |
| xorg | 0.076923076923076 | 0.268141592920354 | 0.048780487804878044 | 0.09037037037037038 |
| jdt | 0.1016949152542373 | 0.2713513513513513 | 0.08993576017130621 | 0.1 |
| lucene | 0.0910031023784901 | 0.26575931232091693 | 0.0563230605738576 | 0.10885167464114832 |

Parameters Used while tuning the models:

1) **Support Vector Machine:** penalty='l1', dual=False, max_iter=10000000, loss='squared_hinge', tol=1e-3

setting the penalty to l1 and dual to false improved the svm model and gave the best value of F1.

2) **Decession Tree:** random_state=None, criterion='gini', max_depth=10, max_leaf_nodes=50, min_samples_leaf=1, splitter='best'

with the above parameters, the classifier gave the best value of F1

3) **Logistic Regression:** random_state=0, max_iter=12000, solver='lbfgs'

4) **Random Forest Regressor:** n_estimators=100, criterion='gini', max_depth=10, min_samples_split=2, min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', min_impurity_decrease=0.0, random_state=None

Part III

| Project | Logistic Regression | Decission Tree | SVM | Random Forest | Naive base model | metrics |
|------------|---------------------|---------------------|---------------------|----------------------|---------------------|-----------|
| jackrabbit | 0.43630166169578183 | 0.4233637116818558 | 0.4455100261551874 | 0.23644251626898047 | 0.4337599446558284 | F1 |
| xorg | 0.15796519410977242 | 0.1611675126903553 | 0.1608725289706885 | 0.026438569206842927 | 0.2725749394898529 | F1 |
| jdt | 0.2702104097452935 | 0.22365339578454335 | 0.21916092673763307 | 0.05309734513274337 | 0.29917032698877505 | F1 |
| lucene | 0.1722113502935421 | 0.14317180616740088 | 0.1366649668536461 | 0.028259473346178548 | 0.31889236266190263 | F1 |

Feature Collection:

To create new set of features for each project, first of all a bag of words was created for each project. The bag of word was created by tokenizing each single patch file and counted the word frequency of all patch files. Then, a dictionary of words of frequency of more than 3 was created of all the tokens collected for the patch files. The dictionary was cleaned of all non words symbols and of tokens of more than 20 chars length.

Then of the tokens of the created dictionary, a loop iterated over all the change_id of the train and test data and calculated the frequency of each key of the dictionary in each change and save it to a numpy array.

The numpy array then converted to a datafram and saved as a new csv file with all the features added to a folder named data.

Comparing results from PartII and PartIII:

By comparing the results of partII and partII, we can see that the F1 vlaues have imporeved for some models while it decreased for others when adding new features and increasing the dimentionality of the input data. The table below shows F1 values for differnt projects using different classifiers that have been trained using partII and partIII data:

| Project | F1/Logistic Reg PartIII | F1/Logistic Reg PartII |
|------------|-------------------------|------------------------|
| jackrabbit | 0.43630166169578183 | 0.38014311270125 |
| xorg | 0.15796519410977242 | 0.076923076923076 |
| jdt | 0.2702104097452935 | 0.1016949152542373 |
| lucene | 0.1722113502935421 | 0.0910031023784901 |

| Project | F1/Decission Tree PartIII | F1/Decission Tree PartII |
|----------------|----------------------------------|---------------------------------|
| jackrabbit | 0.4233637116818558 | 0.3705035971223 |
| xorg | 0.1611675126903553 | 0.268141592920354 |
| jdt | 0.22365339578454335 | 0.2713513513513513 |
| lucene | 0.14317180616740088 | 0.26575931232091693 |

| Project | F1/SVM PartIII | F1/SVM PartII |
|----------------|-----------------------|----------------------|
| jackrabbit | 0.4455100261551874 | 0.3539094650205762 |
| xorg | 0.1608725289706885 | 0.048780487804878044 |
| jdt | 0.21916092673763307 | 0.08993576017130621 |
| lucene | 0.1366649668536461 | 0.0563230605738576 |

| Project | F1/Niave Bays PartIII | F1/Niave Bays PartI |
|----------------|------------------------------|----------------------------|
| jackrabbit | 0.4337599446558284 | 0.2686414708886619 |
| xorg | 0.2725749394898529 | 0.3061667183142237 |
| jdt | 0.29917032698877505 | 0.2731565564838358 |
| lucene | 0.31889236266190263 | 0.35405405405405405 |

Part IV- Improving the training set by resampling

Resampling Methods used:

1) SMOTE

first selects instances from the minority class and finds k nearest neighbors for each instance, where k is a given number. It then creates new instances using the selected instances and their neighbors.

2) Resampling with Replacement

The most common way to address this issue of imbalanced learning is based on a resampling procedure. This approach is simple but it has its own drawbacks. We could decide to upsample the class 1, so as to match the number of samples belonging to class 0. However, we can only use the existing data and, after every sampling step, we restart from the original dataset (replacement).

3) Spread Subsample

This method eliminates instances in the majority class until the ratio of majority instances over minority instances is equal to a given ratio. Weight is a property of the instance that reflects the level of impact this instance has on its class. Spread subsample also updates the weight of each instance in order to maintain the overall weight of the two classes.

| Project | F1/SVM | F1/Decission Tree | Data Sampling method |
|------------|---------------------|---------------------|-----------------------------|
| jackrabbit | 0.5227963525835866 | 0.48393854748603354 | SMOTE |
| xorg | 0.2431736218444101 | 0.3466780238500852 | SMOTE |
| jdt | 0.2778884462151394 | 0.26867119301648884 | SMOTE |
| lucene | 0.2651814588573482 | 0.269539501478665 | SMOTE |
| jackrabbit | 0.5418559377027904 | 0.5036080516521078 | Resampling with Replacement |
| xorg | 0.24489795918367346 | 0.34895833333333337 | Resampling with Replacement |
| jdt | 0.28169014084507044 | 0.32724107919930373 | Resampling with Replacement |
| lucene | 0.22735042735042732 | 0.27761065749892566 | Resampling with Replacement |
| jackrabbit | 0.5835847917923959 | 0.3978269954032595 | Spread Subsample |
| xorg | 0.40176470588235297 | 0.27598566308243727 | Spread Subsample |
| jdt | 0.39955522609340255 | 0.3106976744186047 | Spread Subsample |
| lucene | 0.39270687237026647 | 0.3288186606471031 | Spread Subsample |

For resampling, SVM and Decision tree models were used for the resampled data. By comparing the results that were obtained in partIII and partIV, we notice that with resampling the values of F1 have improved as shown in the table below which show the comparison of results obtained in partIII and partIV.

| Project | F1/ SVM with SMOTE resampling | F1/SVM PartIII |
|----------------|--------------------------------------|-----------------------|
| jackrabbit | 0.5227963525835866 | 0.43630166169578183 |
| xorg | 0.2431736218444101 | 0.15796519410977242 |
| jdt | 0.2778884462151394 | 0.2702104097452935 |
| lucene | 0.2651814588573482 | 0.1722113502935421 |

| Project | F1/Decision Tree with SMOTE resampling | F1/Decision Tree PartIII |
|----------------|---|---------------------------------|
| jackrabbit | 0.48393854748603354 | 0.4233637116818558 |
| xorg | 0.3466780238500852 | 0.1611675126903553 |
| jdt | 0.26867119301648884 | 0.22365339578454335 |
| lucene | 0.269539501478665 | 0.14317180616740088 |

Part V: Further Improvement(bonus)

In this part I used Neural Network deep learning technique. The model that was used to make the prediction is Keras model.

The Keras models were defined as a sequence of layers.

The model was run on the original features given in the data to avoid time and space complexity

A sequential model was created and layers were added one at a time until network architecture is complete to the output layer.

In this model, I used a fully-connected network structure with three layers.

Fully connected layers are defined using the Dense class. The number of neurons or nodes were specified in the layer as the first argument, and the activation function is specified using the **activation** argument.

I used the rectified linear unit activate function known as ReLU on the first two layers and the Sigmoid function in the output layer.

Compiling the model uses the efficient numerical libraries of TensorFlow. The backend automatically chooses the best way to represent the network for training and making predictions.

When compiling, we must specify some additional properties required when training the network. To train the network, it means finding the best set of weights to map inputs to outputs in our dataset.

We must specify the loss function to use to evaluate a set of weights, the optimizer is used to search through different weights for the network and any optional metrics we would like to collect and report during training.

In our case, we will use cross entropy as the **loss** argument. This loss is for a binary classification problems.

Training occurs over epochs and each epoch is split into batches.

- **Epoch:** One pass through all of the rows in the training dataset.
- **Batch:** One or more samples considered by the model within an epoch before weights are updated.

| Project | F1/NN- Keras Model | F1/Decission Tree PartIV(SMOTE) |
|------------|---------------------|------------------------------------|
| jackrabbit | 0.23838162930563148 | 0.48393854748603354 |
| xorg | 0.1816920943134535 | 0.3466780238500852 |
| jdt | 0.073706591070163 | 0.26867119301648884 |
| lucene | 0.16161616161616163 | 0.269539501478665 |

References:

https://scikit-learn.org/stable/supervised_learning.html#supervised-learning

<https://machinelearningmastery.com/tutorial-first-neural-network-python-keras/>